

Complete MySQL Notes - Employee Management System

Table of Contents

1. Introduction to MySQL
 2. Database Design
 3. DDL (Data Definition Language)
 4. DML (Data Manipulation Language)
 5. DQL (Data Query Language)
 6. Advanced Concepts
 7. Database Security
 8. Performance Optimization
-

Introduction to MySQL

MySQL is a popular open-source relational database management system (RDBMS) that uses Structured Query Language (SQL) for managing data.

Key Features:

- **Relational Database:** Data stored in tables with relationships
 - **ACID Properties:** Atomicity, Consistency, Isolation, Durability
 - **Cross-platform:** Works on Windows, Linux, macOS
 - **Scalable:** Handles small to large applications
 - **Open Source:** Free to use with commercial support available
-

Database Design

Employee Management System Schema

Our example will use an employee management system with the following entities:

```
sql
```

```
-- Database Structure Overview
```

```
Departments (dept_id, dept_name, location, manager_id)
```

```
Employees (emp_id, first_name, last_name, email, phone, hire_date, salary, dept_id, manager_id)
```

```
Projects (project_id, project_name, start_date, end_date, budget, dept_id)
```

```
Employee_Projects (emp_id, project_id, role, hours_worked)
```

```
Salaries (emp_id, salary, effective_date)
```



DDL (Data Definition Language)

DDL commands are used to define and modify database structure.

1. CREATE DATABASE

```
sql
```

```
-- Create the database
```

```
CREATE DATABASE employee_management;
```

```
-- Use the database
```

```
USE employee_management;
```

```
-- Show existing databases
```

```
SHOW DATABASES;
```

2. CREATE TABLE

sql

-- Create Departments table

```
CREATE TABLE departments (  
    dept_id INT AUTO_INCREMENT PRIMARY KEY,  
    dept_name VARCHAR(50) NOT NULL UNIQUE,  
    location VARCHAR(100),  
    manager_id INT,  
    created_date TIMESTAMP DEFAULT CURRENT_TIMESTAMP  
);
```

-- Create Employees table

```
CREATE TABLE employees (  
    emp_id INT AUTO_INCREMENT PRIMARY KEY,  
    first_name VARCHAR(50) NOT NULL,  
    last_name VARCHAR(50) NOT NULL,  
    email VARCHAR(100) UNIQUE NOT NULL,  
    phone VARCHAR(15),  
    hire_date DATE NOT NULL,  
    salary DECIMAL(10,2) NOT NULL,  
    dept_id INT,  
    manager_id INT,  
    status ENUM('Active', 'Inactive', 'Terminated') DEFAULT 'Active',  
    created_date TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
    updated_date TIMESTAMP DEFAULT CURRENT_TIMESTAMP ON UPDATE CURRENT_TIMESTAMP,
```

-- Foreign key constraints

```
    FOREIGN KEY (dept_id) REFERENCES departments(dept_id),  
    FOREIGN KEY (manager_id) REFERENCES employees(emp_id)  
);
```

-- Create Projects table

```
CREATE TABLE projects (  
    project_id INT AUTO_INCREMENT PRIMARY KEY,
```

```

project_name VARCHAR(100) NOT NULL,
description TEXT,
start_date DATE NOT NULL,
end_date DATE,
budget DECIMAL(12,2),
status ENUM('Planning', 'Active', 'Completed', 'On Hold') DEFAULT 'Planning',
dept_id INT,
created_date TIMESTAMP DEFAULT CURRENT_TIMESTAMP,

FOREIGN KEY (dept_id) REFERENCES departments(dept_id)
);

-- Create Employee_Projects junction table (Many-to-Many relationship)
CREATE TABLE employee_projects (
    emp_id INT,
    project_id INT,
    role VARCHAR(50),
    hours_worked DECIMAL(5,2) DEFAULT 0,
    assigned_date DATE DEFAULT (CURRENT_DATE),

    PRIMARY KEY (emp_id, project_id),
    FOREIGN KEY (emp_id) REFERENCES employees(emp_id) ON DELETE CASCADE,
    FOREIGN KEY (project_id) REFERENCES projects(project_id) ON DELETE CASCADE
);

-- Create Salary History table
CREATE TABLE salary_history (
    history_id INT AUTO_INCREMENT PRIMARY KEY,
    emp_id INT,
    old_salary DECIMAL(10,2),
    new_salary DECIMAL(10,2),
    change_date DATE DEFAULT (CURRENT_DATE),

```

```
reason VARCHAR(200),  
  
FOREIGN KEY (emp_id) REFERENCES employees(emp_id)  
);
```

3. ALTER TABLE

```
sql  
  
-- Add new column  
ALTER TABLE employees  
ADD COLUMN birth_date DATE;  
  
-- Modify column  
ALTER TABLE employees  
MODIFY COLUMN phone VARCHAR(20);  
  
-- Drop column  
ALTER TABLE employees  
DROP COLUMN birth_date;  
  
-- Add index  
ALTER TABLE employees  
ADD INDEX idx_email (email);  
  
-- Add foreign key constraint  
ALTER TABLE departments  
ADD CONSTRAINT fk_dept_manager  
FOREIGN KEY (manager_id) REFERENCES employees(emp_id);
```

4. DROP and TRUNCATE

```
sql

-- Drop table (removes structure and data)
DROP TABLE IF EXISTS temp_table;

-- Truncate table (removes all data, keeps structure)
TRUNCATE TABLE salary_history;

-- Drop database
DROP DATABASE IF EXISTS old_database;
```

5. CREATE INDEX

```
sql

-- Create single column index
CREATE INDEX idx_last_name ON employees(last_name);

-- Create composite index
CREATE INDEX idx_dept_salary ON employees(dept_id, salary);

-- Create unique index
CREATE UNIQUE INDEX idx_emp_email ON employees(email);

-- Show indexes
SHOW INDEX FROM employees;
```

DML (Data Manipulation Language)

DML commands are used to manipulate data within tables.

1. INSERT

sql

-- Insert into departments

```
INSERT INTO departments (dept_name, location) VALUES
('Human Resources', 'New York'),
('Information Technology', 'San Francisco'),
('Finance', 'Chicago'),
('Marketing', 'Los Angeles'),
('Operations', 'Denver');
```

-- Insert into employees

```
INSERT INTO employees (first_name, last_name, email, phone, hire_date, salary, dept_id) VALUES
('John', 'Smith', 'john.smith@company.com', '555-0101', '2023-01-15', 75000.00, 2),
('Sarah', 'Johnson', 'sarah.johnson@company.com', '555-0102', '2023-02-01', 85000.00, 2),
('Michael', 'Brown', 'michael.brown@company.com', '555-0103', '2023-01-20', 65000.00, 1),
('Emily', 'Davis', 'emily.davis@company.com', '555-0104', '2023-03-10', 70000.00, 3),
('David', 'Wilson', 'david.wilson@company.com', '555-0105', '2023-02-15', 90000.00, 2);
```

-- Update manager_id after employees are inserted

```
UPDATE departments SET manager_id = 3 WHERE dept_id = 1; -- Michael manages HR
UPDATE departments SET manager_id = 5 WHERE dept_id = 2; -- David manages IT
```

-- Insert projects

```
INSERT INTO projects (project_name, description, start_date, end_date, budget, dept_id) VALUES
('Employee Portal', 'Internal employee management system', '2024-01-01', '2024-06-30', 150000.00, 1),
('Payroll System', 'Automated payroll processing', '2024-02-01', '2024-08-31', 200000.00, 2),
('Marketing Campaign', 'Q2 product launch campaign', '2024-03-01', '2024-05-31', 75000.00, 3);
```

-- Insert employee-project assignments

```
INSERT INTO employee_projects (emp_id, project_id, role, hours_worked) VALUES
(1, 1, 'Developer', 120.5),
(2, 1, 'Senior Developer', 150.0),
(5, 1, 'Project Manager', 80.0),
```

```
(4, 2, 'Financial Analyst', 100.0),  
(2, 3, 'Technical Consultant', 40.0);
```

2. UPDATE

sql

-- Simple update

```
UPDATE employees
SET salary = 80000.00
WHERE emp_id = 1;
```

-- Update with calculation

```
UPDATE employees
SET salary = salary * 1.05
WHERE dept_id = 2;
```

-- Update with JOIN

```
UPDATE employees e
JOIN departments d ON e.dept_id = d.dept_id
SET e.salary = e.salary * 1.03
WHERE d.dept_name = 'Finance';
```

-- Update multiple columns

```
UPDATE employees
SET salary = 95000.00, status = 'Active'
WHERE emp_id = 5;
```

-- Record salary change in history

```
INSERT INTO salary_history (emp_id, old_salary, new_salary, reason)
SELECT emp_id, 75000.00, salary, 'Annual Review'
FROM employees
WHERE emp_id = 1;
```

3. DELETE

sql

-- Delete specific record

```
DELETE FROM employee_projects
WHERE emp_id = 2 AND project_id = 3;
```

-- Delete with condition

```
DELETE FROM employees
WHERE status = 'Terminated' AND hire_date < '2020-01-01';
```

-- Delete with subquery

```
DELETE FROM salary_history
WHERE emp_id NOT IN (SELECT emp_id FROM employees);
```

-- Safe delete with EXISTS

```
DELETE e FROM employees e
WHERE EXISTS (
    SELECT 1 FROM employees
    WHERE manager_id = e.emp_id
    AND status = 'Terminated'
);
```

DQL (Data Query Language)

DQL is used to retrieve data from the database.

1. Basic SELECT

sql

-- Select all columns

```
SELECT * FROM employees;
```

-- Select specific columns

```
SELECT first_name, last_name, email, salary FROM employees;
```

-- Select with alias

```
SELECT
    first_name AS 'First Name',
    last_name AS 'Last Name',
    salary AS 'Annual Salary'
FROM employees;
```

-- Select with expressions

```
SELECT
    first_name,
    last_name,
    salary,
    salary * 12 AS annual_salary,
    YEAR(hire_date) AS hire_year
FROM employees;
```

2. WHERE Clause

sql

-- Basic conditions

```
SELECT * FROM employees WHERE salary > 70000;
```

-- Multiple conditions

```
SELECT * FROM employees  
WHERE dept_id = 2 AND salary BETWEEN 70000 AND 90000;
```

-- Pattern matching

```
SELECT * FROM employees  
WHERE first_name LIKE 'J%' OR last_name LIKE '%son';
```

-- IN operator

```
SELECT * FROM employees  
WHERE dept_id IN (1, 2, 3);
```

-- NULL checks

```
SELECT * FROM projects  
WHERE end_date IS NULL;
```

-- Date conditions

```
SELECT * FROM employees  
WHERE hire_date >= '2023-01-01' AND hire_date < '2024-01-01';
```

3. ORDER BY and LIMIT

sql

-- Single column sorting

```
SELECT * FROM employees ORDER BY salary DESC;
```

-- Multiple column sorting

```
SELECT * FROM employees  
ORDER BY dept_id ASC, salary DESC;
```

-- Limit results

```
SELECT * FROM employees  
ORDER BY salary DESC  
LIMIT 5;
```

-- Pagination

```
SELECT * FROM employees  
ORDER BY emp_id  
LIMIT 10 OFFSET 20;
```

4. Aggregate Functions

sql

-- Count employees

```
SELECT COUNT(*) AS total_employees FROM employees;
```

-- Count by department

```
SELECT
    d.dept_name,
    COUNT(e.emp_id) AS employee_count
FROM departments d
LEFT JOIN employees e ON d.dept_id = e.dept_id
GROUP BY d.dept_id, d.dept_name;
```

-- Salary statistics

```
SELECT
    AVG(salary) AS avg_salary,
    MIN(salary) AS min_salary,
    MAX(salary) AS max_salary,
    SUM(salary) AS total_payroll
FROM employees;
```

-- Department-wise statistics

```
SELECT
    d.dept_name,
    COUNT(e.emp_id) AS emp_count,
    AVG(e.salary) AS avg_salary,
    SUM(e.salary) AS total_payroll
FROM departments d
LEFT JOIN employees e ON d.dept_id = e.dept_id
GROUP BY d.dept_id, d.dept_name
HAVING COUNT(e.emp_id) > 0
ORDER BY avg_salary DESC;
```

5. GROUP BY and HAVING

sql

-- Group by single column

```
SELECT dept_id, COUNT(*) AS emp_count, AVG(salary) AS avg_salary
FROM employees
GROUP BY dept_id;
```

-- Group by multiple columns

```
SELECT
    YEAR(hire_date) AS hire_year,
    dept_id,
    COUNT(*) AS new_hires
FROM employees
GROUP BY YEAR(hire_date), dept_id
ORDER BY hire_year, dept_id;
```

-- HAVING clause

```
SELECT
    dept_id,
    AVG(salary) AS avg_salary
FROM employees
GROUP BY dept_id
HAVING AVG(salary) > 70000;
```

6. JOINS

sql

```
-- INNER JOIN
SELECT
    e.first_name,
    e.last_name,
    d.dept_name,
    e.salary
FROM employees e
INNER JOIN departments d ON e.dept_id = d.dept_id;
```

```
-- LEFT JOIN
SELECT
    d.dept_name,
    e.first_name,
    e.last_name
FROM departments d
LEFT JOIN employees e ON d.dept_id = e.dept_id;
```

```
-- RIGHT JOIN
SELECT
    e.first_name,
    e.last_name,
    d.dept_name
FROM employees e
RIGHT JOIN departments d ON e.dept_id = d.dept_id;
```

```
-- FULL OUTER JOIN (MySQL doesn't support directly, use UNION)
SELECT d.dept_name, e.first_name, e.last_name
FROM departments d
LEFT JOIN employees e ON d.dept_id = e.dept_id
UNION
SELECT d.dept_name, e.first_name, e.last_name
FROM departments d
```

```
RIGHT JOIN employees e ON d.dept_id = e.dept_id;
```

```
-- Self JOIN (Employee and Manager)
```

```
SELECT
```

```
    e.first_name AS employee_name,
```

```
    m.first_name AS manager_name
```

```
FROM employees e
```

```
LEFT JOIN employees m ON e.manager_id = m.emp_id;
```

```
-- Multiple JOINS
```

```
SELECT
```

```
    e.first_name,
```

```
    e.last_name,
```

```
    d.dept_name,
```

```
    p.project_name,
```

```
    ep.role,
```

```
    ep.hours_worked
```

```
FROM employees e
```

```
JOIN departments d ON e.dept_id = d.dept_id
```

```
JOIN employee_projects ep ON e.emp_id = ep.emp_id
```

```
JOIN projects p ON ep.project_id = p.project_id;
```

7. Subqueries

sql

-- Scalar subquery

```
SELECT first_name, last_name, salary
FROM employees
WHERE salary > (SELECT AVG(salary) FROM employees);
```

-- Correlated subquery

```
SELECT e1.first_name, e1.last_name, e1.salary, e1.dept_id
FROM employees e1
WHERE e1.salary > (
    SELECT AVG(e2.salary)
    FROM employees e2
    WHERE e2.dept_id = e1.dept_id
);
```

-- EXISTS subquery

```
SELECT d.dept_name
FROM departments d
WHERE EXISTS (
    SELECT 1 FROM employees e
    WHERE e.dept_id = d.dept_id
);
```

-- IN subquery

```
SELECT first_name, last_name
FROM employees
WHERE dept_id IN (
    SELECT dept_id FROM departments
    WHERE location = 'San Francisco'
);
```

-- Subquery in FROM clause

```
SELECT dept_summary.dept_name, dept_summary.avg_salary
```



```
FROM (  
    SELECT d.dept_name, AVG(e.salary) as avg_salary  
    FROM departments d  
    JOIN employees e ON d.dept_id = e.dept_id  
    GROUP BY d.dept_id, d.dept_name  
) AS dept_summary  
WHERE dept_summary.avg_salary > 75000;
```

Advanced Concepts

1. Views

sql

-- Create a view for employee details

CREATE VIEW employee_details AS

SELECT

 e.emp_id,
 CONCAT(e.first_name, ' ', e.last_name) AS full_name,
 e.email,
 e.salary,
 d.dept_name,
 CONCAT(m.first_name, ' ', m.last_name) AS manager_name

FROM employees e

LEFT JOIN departments d ON e.dept_id = d.dept_id

LEFT JOIN employees m ON e.manager_id = m.emp_id

WHERE e.status = 'Active';

-- Use the view

SELECT * FROM employee_details WHERE salary > 70000;

-- Create view for department summary

CREATE VIEW department_summary AS

SELECT

 d.dept_name,
 d.location,
 COUNT(e.emp_id) AS employee_count,
 AVG(e.salary) AS avg_salary,
 SUM(e.salary) AS total_payroll

FROM departments d

LEFT JOIN employees e ON d.dept_id = e.dept_id AND e.status = 'Active'

GROUP BY d.dept_id, d.dept_name, d.location;

-- Drop view

DROP VIEW IF EXISTS employee_details;

2. Stored Procedures

sql

```

-- Simple stored procedure
DELIMITER //
CREATE PROCEDURE GetEmployeesByDepartment(IN dept_name VARCHAR(50))
BEGIN
    SELECT e.first_name, e.last_name, e.salary
    FROM employees e
    JOIN departments d ON e.dept_id = d.dept_id
    WHERE d.dept_name = dept_name AND e.status = 'Active';
END //
DELIMITER ;

-- Call the procedure
CALL GetEmployeesByDepartment('Information Technology');

-- Procedure with OUT parameter
DELIMITER //
CREATE PROCEDURE GetDepartmentStats(
    IN dept_name VARCHAR(50),
    OUT emp_count INT,
    OUT avg_salary DECIMAL(10,2)
)
BEGIN
    SELECT COUNT(*), AVG(salary)
    INTO emp_count, avg_salary
    FROM employees e
    JOIN departments d ON e.dept_id = d.dept_id
    WHERE d.dept_name = dept_name AND e.status = 'Active';
END //
DELIMITER ;

-- Call procedure with OUT parameters
CALL GetDepartmentStats('Information Technology', @count, @avg_sal);

```

```

SELECT @count AS employee_count, @avg_sal AS average_salary;

-- Procedure for salary increase
DELIMITER //
CREATE PROCEDURE GiveSalaryIncrease(
    IN emp_id INT,
    IN increase_percent DECIMAL(5,2),
    OUT new_salary DECIMAL(10,2)
)
BEGIN
    DECLARE old_salary DECIMAL(10,2);

    -- Get current salary
    SELECT salary INTO old_salary FROM employees WHERE employees.emp_id = emp_id;

    -- Calculate new salary
    SET new_salary = old_salary * (1 + increase_percent / 100);

    -- Update employee salary
    UPDATE employees SET salary = new_salary WHERE employees.emp_id = emp_id;

    -- Log the change
    INSERT INTO salary_history (emp_id, old_salary, new_salary, reason)
    VALUES (emp_id, old_salary, new_salary, CONCAT('Salary increase of ', increase_percent));
END //
DELIMITER ;

```

3. Functions

sql


```
-- Create a function to calculate experience
DELIMITER //
CREATE FUNCTION CalculateExperience(hire_date DATE)
RETURNS DECIMAL(4,2)
READS SQL DATA
DETERMINISTIC
BEGIN
    DECLARE experience DECIMAL(4,2);
    SET experience = TIMESTAMPDIFF(MONTH, hire_date, CURDATE()) / 12;
    RETURN experience;
END //
DELIMITER ;
```

```
-- Use the function
SELECT
    first_name,
    last_name,
    hire_date,
    CalculateExperience(hire_date) AS years_experience
FROM employees;
```

```
-- Function to get employee grade based on salary
DELIMITER //
CREATE FUNCTION GetEmployeeGrade(salary DECIMAL(10,2))
RETURNS VARCHAR(10)
DETERMINISTIC
BEGIN
    DECLARE grade VARCHAR(10);

    IF salary >= 90000 THEN
        SET grade = 'Senior';
    ELSEIF salary >= 70000 THEN
```

```
        SET grade = 'Mid-level';
    ELSE
        SET grade = 'Junior';
    END IF;

    RETURN grade;
END //
DELIMITER ;

-- Use the grade function
SELECT
    first_name,
    last_name,
    salary,
    GetEmployeeGrade(salary) AS employee_grade
FROM employees;
```

4. Triggers

sql

```

-- Trigger to automatically update salary history
DELIMITER //
CREATE TRIGGER salary_change_trigger
AFTER UPDATE ON employees
FOR EACH ROW
BEGIN
    IF OLD.salary != NEW.salary THEN
        INSERT INTO salary_history (emp_id, old_salary, new_salary, reason)
        VALUES (NEW.emp_id, OLD.salary, NEW.salary, 'Salary Updated');
    END IF;
END //
DELIMITER ;

-- Trigger to prevent deletion of employees with active projects
DELIMITER //
CREATE TRIGGER prevent_employee_deletion
BEFORE DELETE ON employees
FOR EACH ROW
BEGIN
    DECLARE project_count INT;

    SELECT COUNT(*) INTO project_count
    FROM employee_projects ep
    JOIN projects p ON ep.project_id = p.project_id
    WHERE ep.emp_id = OLD.emp_id AND p.status = 'Active';

    IF project_count > 0 THEN
        SIGNAL SQLSTATE '45000'
        SET MESSAGE_TEXT = 'Cannot delete employee with active projects';
    END IF;
END //
DELIMITER ;

```

```
-- Trigger to update project budget when employee hours change
DELIMITER //
CREATE TRIGGER update_project_hours
AFTER UPDATE ON employee_projects
FOR EACH ROW
BEGIN
    DECLARE total_hours DECIMAL(8,2);

    IF OLD.hours_worked != NEW.hours_worked THEN
        SELECT SUM(hours_worked) INTO total_hours
        FROM employee_projects
        WHERE project_id = NEW.project_id;

        -- You could add logic here to update project status or costs
        -- based on total hours worked
    END IF;
END //
DELIMITER ;
```

5. Transactions

sql

```
-- Basic transaction example
START TRANSACTION;

INSERT INTO employees (first_name, last_name, email, hire_date, salary, dept_id)
VALUES ('Alice', 'Cooper', 'alice.cooper@company.com', '2024-01-15', 72000.00, 2);

SET @new_emp_id = LAST_INSERT_ID();

INSERT INTO employee_projects (emp_id, project_id, role)
VALUES (@new_emp_id, 1, 'Developer');

COMMIT;

-- Transaction with rollback
START TRANSACTION;

UPDATE employees SET salary = salary * 1.10 WHERE dept_id = 2;

-- Check if update is acceptable
SELECT COUNT(*) as high_salary_count
FROM employees
WHERE salary > 100000 AND dept_id = 2;

-- If too many high salaries, rollback
ROLLBACK;

-- Transaction with savepoints
START TRANSACTION;

SAVEPOINT before_salary_update;

UPDATE employees SET salary = salary * 1.05 WHERE dept_id = 1;
```

```
SAVEPOINT after_hr_update;
```

```
UPDATE employees SET salary = salary * 1.08 WHERE dept_id = 2;
```

```
-- If something goes wrong, can rollback to savepoint
```

```
ROLLBACK TO before_salary_update;
```

```
COMMIT;
```

6. Common Table Expressions (CTEs)

sql

```

-- Simple CTE
WITH department_avg AS (
    SELECT dept_id, AVG(salary) as avg_salary
    FROM employees
    GROUP BY dept_id
)
SELECT e.first_name, e.last_name, e.salary, da.avg_salary
FROM employees e
JOIN department_avg da ON e.dept_id = da.dept_id
WHERE e.salary > da.avg_salary;

-- Recursive CTE for organizational hierarchy
WITH RECURSIVE employee_hierarchy AS (
    -- Base case: top-level employees (no manager)
    SELECT emp_id, first_name, last_name, manager_id, 0 as level
    FROM employees
    WHERE manager_id IS NULL

    UNION ALL

    -- Recursive case: employees with managers
    SELECT e.emp_id, e.first_name, e.last_name, e.manager_id, eh.level + 1
    FROM employees e
    JOIN employee_hierarchy eh ON e.manager_id = eh.emp_id
)
SELECT emp_id,
       CONCAT(REPEAT(' ', level), first_name, ' ', last_name) as employee_hierarchy,
       level
FROM employee_hierarchy
ORDER BY level, last_name;

-- Multiple CTEs

```

```
WITH
high_performers AS (
    SELECT emp_id, first_name, last_name, salary
    FROM employees
    WHERE salary > (SELECT AVG(salary) FROM employees)
),
project_leaders AS (
    SELECT DISTINCT ep.emp_id
    FROM employee_projects ep
    WHERE ep.role LIKE '%Manager%' OR ep.role LIKE '%Lead%'
)
SELECT hp.first_name, hp.last_name, hp.salary
FROM high_performers hp
JOIN project_leaders pl ON hp.emp_id = pl.emp_id;
```

Database Security

1. User Management

sql

-- Create users

```
CREATE USER 'hr_manager'@'localhost' IDENTIFIED BY 'secure_password';  
CREATE USER 'developer'@'%' IDENTIFIED BY 'dev_password';  
CREATE USER 'readonly_user'@'localhost' IDENTIFIED BY 'read_password';
```

-- Grant privileges

```
GRANT ALL PRIVILEGES ON employee_management.* TO 'hr_manager'@'localhost';  
GRANT SELECT, INSERT, UPDATE ON employee_management.employees TO 'developer'@'%';  
GRANT SELECT ON employee_management.* TO 'readonly_user'@'localhost';
```

-- Grant specific column privileges

```
GRANT SELECT (first_name, last_name, email) ON employees TO 'limited_user'@'localhost';
```

-- Revoke privileges

```
REVOKE INSERT ON employee_management.employees FROM 'developer'@'%';
```

-- Show privileges

```
SHOW GRANTS FOR 'hr_manager'@'localhost';
```

-- Drop user

```
DROP USER 'old_user'@'localhost';
```

2. Data Validation and Constraints

sql

-- Add check constraints (MySQL 8.0+)

ALTER TABLE employees

ADD CONSTRAINT chk_salary CHECK (salary > 0);

ALTER TABLE employees

ADD CONSTRAINT chk_hire_date CHECK (hire_date <= CURDATE());

-- Create table with constraints

CREATE TABLE employee_reviews (

review_id INT AUTO_INCREMENT PRIMARY KEY,

emp_id INT,

review_date DATE DEFAULT (CURDATE()),

rating DECIMAL(3,2),

comments TEXT,

CONSTRAINT fk_review_emp FOREIGN KEY (emp_id) REFERENCES employees(emp_id),

CONSTRAINT chk_rating CHECK (rating >= 1.0 AND rating <= 5.0),

CONSTRAINT chk_review_date CHECK (review_date <= CURDATE())

);

Performance Optimization

1. Indexing Strategies

sql

-- Analyze query performance

```
EXPLAIN SELECT * FROM employees WHERE last_name = 'Smith';
```

-- Create covering index

```
CREATE INDEX idx_emp_dept_salary ON employees(dept_id, salary, first_name, last_name);
```

-- Create partial index (MySQL doesn't support, but here's the concept)

```
-- CREATE INDEX idx_active_employees ON employees(dept_id) WHERE status = 'Active';
```

-- Monitor index usage

```
SELECT
```

```
    table_name,
```

```
    index_name,
```

```
    column_name,
```

```
    cardinality
```

```
FROM information_schema.statistics
```

```
WHERE table_schema = 'employee_management'
```

```
ORDER BY table_name, index_name;
```

2. Query Optimization

sql

-- Inefficient query

```
SELECT * FROM employees e1
WHERE salary > (
    SELECT AVG(salary) FROM employees e2 WHERE e2.dept_id = e1.dept_id
);
```

-- Optimized version using JOIN

```
SELECT e.*, dept_avg.avg_salary
FROM employees e
JOIN (
    SELECT dept_id, AVG(salary) as avg_salary
    FROM employees
    GROUP BY dept_id
) dept_avg ON e.dept_id = dept_avg.dept_id
WHERE e.salary > dept_avg.avg_salary;
```

-- Use LIMIT for large result sets

```
SELECT * FROM employees
ORDER BY hire_date DESC
LIMIT 10;
```

*-- Avoid SELECT * when possible*

```
SELECT emp_id, first_name, last_name, salary
FROM employees
WHERE dept_id = 2;
```

3. Database Maintenance

sql

-- Analyze table statistics

ANALYZE TABLE employees;

-- Optimize table

OPTIMIZE TABLE employees;

-- Check table integrity

CHECK TABLE employees;

-- Repair table if needed

REPAIR TABLE employees;

-- Show table status

SHOW TABLE STATUS FROM employee_management;

Practical Examples and Use Cases

1. Employee Reporting Queries

sql

-- Monthly hiring report

```
SELECT
    DATE_FORMAT(hire_date, '%Y-%m') as hire_month,
    COUNT(*) as new_hires,
    AVG(salary) as avg_starting_salary
FROM employees
WHERE hire_date >= DATE_SUB(CURDATE(), INTERVAL 12 MONTH)
GROUP BY DATE_FORMAT(hire_date, '%Y-%m')
ORDER BY hire_month;
```

-- Department performance dashboard

```
SELECT
    d.dept_name,
    COUNT(e.emp_id) as employee_count,
    AVG(e.salary) as avg_salary,
    SUM(CASE WHEN e.hire_date >= DATE_SUB(CURDATE(), INTERVAL 1 YEAR) THEN 1 ELSE 0 END) ;
    COUNT(p.project_id) as active_projects
FROM departments d
LEFT JOIN employees e ON d.dept_id = e.dept_id AND e.status = 'Active'
LEFT JOIN projects p ON d.dept_id = p.dept_id AND p.status = 'Active'
GROUP BY d.dept_id, d.dept_name
ORDER BY employee_count DESC;
```

-- Employee project workload

```
SELECT
    e.first_name,
    e.last_name,
    COUNT(ep.project_id) as project_count,
    SUM(ep.hours_worked) as total_hours,
    AVG(ep.hours_worked) as avg_hours_per_project
FROM employees e
JOIN employee_projects ep ON e.emp_id = ep.emp_id
```

```
JOIN projects p ON ep.project_id = p.project_id
WHERE p.status IN ('Active', 'Planning')
GROUP BY e.emp_id, e.first_name, e.last_name
HAVING project_count > 1
ORDER BY total_hours DESC;
```

2. Data Migration and Cleanup

```

sql

-- Archive old salary history
CREATE TABLE salary_history_archive AS
SELECT * FROM salary_history
WHERE change_date < DATE_SUB(CURDATE(), INTERVAL 2 YEAR);

DELETE FROM salary_history
WHERE change_date < DATE_SUB(CURDATE(), INTERVAL 2 YEAR);

-- Update employee status based on project activity
UPDATE employees
SET status = 'Inactive'
WHERE emp_id NOT IN (
    SELECT DISTINCT ep.emp_id
    FROM employee_projects ep
    JOIN projects p ON ep.project_id = p.project_id
    WHERE p.status = 'Active'
)
AND status = 'Active'
AND hire_date < DATE_SUB(CURDATE(), INTERVAL 6 MONTH);

-- Clean up duplicate email addresses
DELETE e1 FROM employees e1
INNER JOIN employees e2
WHERE e1.emp_id < e2.emp_id
AND e1.email = e2.email;

```

3. Advanced Analytics Queries

sql

-- Salary distribution analysis

```
SELECT
    CASE
        WHEN salary < 50000 THEN 'Under 50K'
        WHEN salary BETWEEN 50000 AND 70000 THEN '50K-70K'
        WHEN salary BETWEEN 70001 AND 90000 THEN '70K-90K'
        WHEN salary > 90000 THEN 'Over 90K'
    END AS salary_range,
    COUNT(*) AS employee_count,
    ROUND(COUNT(*) * 100.0 / (SELECT COUNT(*) FROM employees), 2) AS percentage
FROM employees
WHERE status = 'Active'
GROUP BY salary_range
ORDER BY MIN(salary);
```

-- Employee retention analysis

```
SELECT
    d.dept_name,
    COUNT(e.emp_id) AS total_employees,
    SUM(CASE WHEN e.hire_date >= DATE_SUB(CURDATE(), INTERVAL 1 YEAR) THEN 1 ELSE 0 END) AS
    SUM(CASE WHEN e.status = 'Terminated' AND e.updated_date >= DATE_SUB(CURDATE(), INTERVAL 1 YEAR) THEN 1 ELSE 0 END) AS
    ROUND(
        (COUNT(e.emp_id) - SUM(CASE WHEN e.status = 'Terminated' AND e.updated_date >= DATE_SUB(CURDATE(), INTERVAL 1 YEAR) THEN 1 ELSE 0 END)) /
        2
    ) AS retention_rate
FROM departments d
LEFT JOIN employees e ON d.dept_id = e.dept_id
GROUP BY d.dept_id, d.dept_name
ORDER BY retention_rate DESC;
```

-- Project timeline and resource allocation

```
SELECT
```

```
p.project_name,  
p.start_date,  
p.end_date,  
DATEDIFF(p.end_date, p.start_date) as duration_days,  
COUNT(ep.emp_id) as team_size,  
SUM(ep.hours_worked) as total_hours,  
p.budget,  
ROUND(p.budget / NULLIF(SUM(ep.hours_worked), 0), 2) as cost_per_hour  
FROM projects p  
LEFT JOIN employee_projects ep ON p.project_id = ep.project_id  
WHERE p.status IN ('Active', 'Completed')  
GROUP BY p.project_id, p.project_name, p.start_date, p.end_date, p.budget  
ORDER BY p.start_date DESC;
```

Window Functions (MySQL 8.0+)

Window functions perform calculations across a set of table rows related to the current row.

1. Ranking Functions

sql

-- Rank employees by salary within each department

```
SELECT
    first_name,
    last_name,
    dept_id,
    salary,
    RANK() OVER (PARTITION BY dept_id ORDER BY salary DESC) as salary_rank,
    DENSE_RANK() OVER (PARTITION BY dept_id ORDER BY salary DESC) as dense_rank,
    ROW_NUMBER() OVER (PARTITION BY dept_id ORDER BY salary DESC) as row_num
FROM employees
WHERE status = 'Active';
```

-- Top 3 highest paid employees per department

```
SELECT *
FROM (
    SELECT
        e.first_name,
        e.last_name,
        d.dept_name,
        e.salary,
        RANK() OVER (PARTITION BY e.dept_id ORDER BY e.salary DESC) as salary_rank
    FROM employees e
    JOIN departments d ON e.dept_id = d.dept_id
    WHERE e.status = 'Active'
) ranked_employees
WHERE salary_rank <= 3;
```

-- Percentile ranking

```
SELECT
    first_name,
    last_name,
    salary,
```

```
PERCENT_RANK() OVER (ORDER BY salary) as percentile_rank,  
CUME_DIST() OVER (ORDER BY salary) as cumulative_distribution  
FROM employees  
WHERE status = 'Active';
```

2. Aggregate Window Functions

sql

-- Running total of salaries

```
SELECT
    first_name,
    last_name,
    hire_date,
    salary,
    SUM(salary) OVER (ORDER BY hire_date) as running_total_payroll
FROM employees
WHERE status = 'Active'
ORDER BY hire_date;
```

-- Moving average salary (3-employee window)

```
SELECT
    first_name,
    last_name,
    hire_date,
    salary,
    AVG(salary) OVER (
        ORDER BY hire_date
        ROWS BETWEEN 2 PRECEDING AND CURRENT ROW
    ) as moving_avg_salary
FROM employees
WHERE status = 'Active'
ORDER BY hire_date;
```

-- Department salary statistics with individual comparison

```
SELECT
    e.first_name,
    e.last_name,
    d.dept_name,
    e.salary,
    AVG(e.salary) OVER (PARTITION BY e.dept_id) as dept_avg_salary,
```

```
MAX(e.salary) OVER (PARTITION BY e.dept_id) as dept_max_salary,  
MIN(e.salary) OVER (PARTITION BY e.dept_id) as dept_min_salary,  
e.salary - AVG(e.salary) OVER (PARTITION BY e.dept_id) as salary_vs_dept_avg  
FROM employees e  
JOIN departments d ON e.dept_id = d.dept_id  
WHERE e.status = 'Active';
```

3. Value Functions

sql

-- Compare with previous and next employee salaries

```
SELECT
    first_name,
    last_name,
    hire_date,
    salary,
    LAG(salary, 1) OVER (ORDER BY hire_date) as prev_hire_salary,
    LEAD(salary, 1) OVER (ORDER BY hire_date) as next_hire_salary,
    salary - LAG(salary, 1) OVER (ORDER BY hire_date) as salary_diff_from_prev
FROM employees
WHERE status = 'Active'
ORDER BY hire_date;
```

-- First and last values in each department

```
SELECT
    first_name,
    last_name,
    dept_id,
    hire_date,
    salary,
    FIRST_VALUE(salary) OVER (
        PARTITION BY dept_id
        ORDER BY hire_date
        ROWS UNBOUNDED PRECEDING
    ) as first_dept_salary,
    LAST_VALUE(salary) OVER (
        PARTITION BY dept_id
        ORDER BY hire_date
        ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING
    ) as last_dept_salary
```

```
FROM employees  
WHERE status = 'Active';
```

JSON Data Type and Functions (MySQL 5.7+)

MySQL supports JSON data type for storing and manipulating JSON documents.

1. JSON Column and Basic Operations

sql

```

-- Add JSON column to employees table
ALTER TABLE employees ADD COLUMN skills JSON;
ALTER TABLE employees ADD COLUMN contact_info JSON;

-- Insert JSON data
UPDATE employees
SET skills = JSON_ARRAY('MySQL', 'Python', 'JavaScript'),
    contact_info = JSON_OBJECT(
        'emergency_contact', JSON_OBJECT('name', 'Jane Smith', 'phone', '555-9999'),
        'address', JSON_OBJECT('street', '123 Main St', 'city', 'New York', 'zip', '10001')
    )
WHERE emp_id = 1;

UPDATE employees
SET skills = JSON_ARRAY('Java', 'Spring', 'React'),
    contact_info = JSON_OBJECT(
        'emergency_contact', JSON_OBJECT('name', 'Bob Johnson', 'phone', '555-8888'),
        'address', JSON_OBJECT('street', '456 Oak Ave', 'city', 'San Francisco', 'zip', '94101')
    )
WHERE emp_id = 2;

-- Query JSON data
SELECT
    first_name,
    last_name,
    JSON_EXTRACT(skills, '$[0]') as primary_skill,
    JSON_EXTRACT(contact_info, '$.emergency_contact.name') as emergency_contact
FROM employees
WHERE skills IS NOT NULL;

-- Using -> and ->> operators (shorthand for JSON_EXTRACT)
SELECT

```

```
    first_name,  
    last_name,  
    skills->'${0}' as primary_skill,  
    contact_info->>'$.address.city' as city  
FROM employees  
WHERE skills IS NOT NULL;
```

2. JSON Functions

sql

-- JSON search and manipulation

```
SELECT
    first_name,
    last_name,
    skills,
    JSON_LENGTH(skills) as skill_count,
    JSON_CONTAINS(skills, '"MySQL"') as knows_mysql,
    JSON_SEARCH(skills, 'one', 'Python') as python_position
FROM employees
WHERE skills IS NOT NULL;
```

-- JSON array operations

```
SELECT
    first_name,
    last_name,
    skills,
    JSON_ARRAY_APPEND(skills, ', 'Docker') as skills_with_docker,
    JSON_ARRAY_INSERT(skills, '$[1]', 'Git') as skills_with_git
FROM employees
WHERE emp_id = 1;
```

-- JSON object operations

```
SELECT
    first_name,
    last_name,
    contact_info,
    JSON_SET(contact_info, '$.phone', '555-1234') as updated_contact,
    JSON_REMOVE(contact_info, '$.address.zip') as contact_without_zip
FROM employees
WHERE contact_info IS NOT NULL
LIMIT 1;
```

3. JSON Indexing

sql

-- Create functional index on JSON column

ALTER TABLE employees

ADD INDEX idx_primary_skill ((CAST(skills->'\$[0]' AS CHAR(50))));

-- Create index on JSON path

ALTER TABLE employees

ADD INDEX idx_city ((CAST(contact_info->'\$.address.city' AS CHAR(50))));

-- Query using JSON index

SELECT first_name, last_name, skills

FROM employees

WHERE skills->'\$[0]' = 'MySQL';

Backup and Recovery

1. Logical Backup with mysqldump

```
bash
```

```
# Complete database backup
```

```
mysqldump -u root -p employee_management > employee_management_backup.sql
```

```
# Backup specific tables
```

```
mysqldump -u root -p employee_management employees departments > partial_backup.sql
```

```
# Backup with additional options
```

```
mysqldump -u root -p \  
  --single-transaction \  
  --routines \  
  --triggers \  
  --events \  
  employee_management > complete_backup.sql
```

```
# Backup all databases
```

```
mysqldump -u root -p --all-databases > all_databases_backup.sql
```

2. Point-in-Time Recovery

```
sql
```

```
-- Enable binary logging (in my.cnf)
-- log-bin=mysql-bin
-- binlog-format=ROW

-- Show binary Logs
SHOW BINARY LOGS;

-- Show binary Log events
SHOW BINLOG EVENTS IN 'mysql-bin.000001';

-- Create a backup point
FLUSH LOGS;
```

3. Database Restoration

```
bash
```

```
# Restore from backup
mysql -u root -p employee_management < employee_management_backup.sql

# Restore specific tables
mysql -u root -p employee_management < partial_backup.sql

# Create database before restore
mysql -u root -p -e "CREATE DATABASE employee_management_restored;"
mysql -u root -p employee_management_restored < employee_management_backup.sql
```

Monitoring and Troubleshooting

1. Performance Monitoring

```
sql

-- Show running processes
SHOW PROCESSLIST;

-- Show slow queries (enable slow query log first)
-- In my.cnf: slow_query_log = 1, long_query_time = 2

-- Check table locks
SHOW OPEN TABLES WHERE In_use > 0;

-- Show engine status
SHOW ENGINE INNODB STATUS;

-- Query cache statistics
SHOW STATUS LIKE 'Qcache%';

-- Connection statistics
SHOW STATUS LIKE 'Connections';
SHOW STATUS LIKE 'Threads_connected';
SHOW STATUS LIKE 'Max_used_connections';
```

2. Space Usage Analysis

sql

-- Database size

SELECT

table_schema AS 'Database',

ROUND(SUM(data_length + index_length) / 1024 / 1024, 2) AS 'Size (MB)'

FROM information_schema.tables

WHERE table_schema = 'employee_management'

GROUP BY table_schema;

-- Table sizes

SELECT

table_name AS 'Table',

ROUND(((data_length + index_length) / 1024 / 1024), 2) AS 'Size (MB)',

table_rows AS 'Rows'

FROM information_schema.tables

WHERE table_schema = 'employee_management'

ORDER BY (data_length + index_length) DESC;

-- Index usage

SELECT

t.table_name,

t.index_name,

t.column_name,

s.cardinality,

ROUND(s.cardinality / tr.table_rows * 100, 2) as selectivity_percent

FROM information_schema.statistics s

JOIN information_schema.tables tr ON s.table_name = tr.table_name

JOIN information_schema.statistics t ON s.table_name = t.table_name

WHERE s.table_schema = 'employee_management'

AND tr.table_schema = 'employee_management'

ORDER BY selectivity_percent DESC;

Best Practices and Tips

1. Database Design Best Practices

sql

-- Use appropriate data types

-- Good

```
CREATE TABLE example_good (  
    id INT AUTO_INCREMENT PRIMARY KEY,  
    status ENUM('active', 'inactive') NOT NULL,  
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,  
    amount DECIMAL(10,2) NOT NULL  
);
```

-- Avoid

```
CREATE TABLE example_bad (  
    id VARCHAR(50) PRIMARY KEY, -- Use INT for IDs  
    status VARCHAR(255),         -- Use ENUM for limited options  
    created_at VARCHAR(50),      -- Use proper date/time types  
    amount FLOAT                 -- Use DECIMAL for money  
);
```

-- Normalize your database

-- Instead of storing repeated department info in employees table

-- employees (emp_id, name, dept_name, dept_location) -- Bad

-- Use separate tables with foreign keys

-- departments (dept_id, dept_name, location) -- Good

-- employees (emp_id, name, dept_id) -- Good

2. Query Optimization Tips

sql

-- Use EXPLAIN to analyze queries

EXPLAIN FORMAT=JSON

SELECT e.first_name, e.last_name, d.dept_name

FROM employees e

JOIN departments d ON e.dept_id = d.dept_id

WHERE e.salary > 70000;

*-- Avoid SELECT * in production*

-- Bad

SELECT * FROM employees WHERE dept_id = 2;

-- Good

SELECT emp_id, first_name, last_name, salary FROM employees WHERE dept_id = 2;

-- Use LIMIT for large result sets

SELECT first_name, last_name, hire_date

FROM employees

ORDER BY hire_date DESC

LIMIT 10;

-- Use EXISTS instead of IN for large subqueries

-- Less efficient

SELECT * FROM employees

WHERE dept_id IN (SELECT dept_id FROM departments WHERE location = 'New York');

-- More efficient

SELECT * FROM employees e

WHERE EXISTS (

SELECT 1 FROM departments d

WHERE d.dept_id = e.dept_id AND d.location = 'New York'

);

3. Security Best Practices

sql

```
-- Use parameterized queries (in application code)
-- Instead of: SELECT * FROM employees WHERE emp_id = '' + user_input + ''
-- Use prepared statements with parameters

-- Principle of least privilege
CREATE USER 'app_user'@'localhost' IDENTIFIED BY 'strong_password';
GRANT SELECT, INSERT, UPDATE ON employee_management.employees TO 'app_user'@'localhost';
-- Don't grant unnecessary privileges

-- Regular security maintenance
-- Update passwords regularly
ALTER USER 'app_user'@'localhost' IDENTIFIED BY 'new_strong_password';

-- Remove unused users
DROP USER 'old_user'@'localhost';

-- Monitor failed login attempts
-- Enable general log to track access
-- SET GLOBAL general_log = 'ON';
```



4. Maintenance Tasks

```
sql

-- Regular maintenance procedures
-- Schedule these tasks during low-usage periods

-- Update table statistics
ANALYZE TABLE employees, departments, projects;

-- Defragment tables
OPTIMIZE TABLE employees;

-- Check for corruption
CHECK TABLE employees EXTENDED;

-- Update configuration for better performance
-- In my.cnf:
-- innodb_buffer_pool_size = 70% of available RAM
-- query_cache_size = 256M (if using MySQL < 8.0)
-- max_connections = appropriate for your application
```

Common Interview Questions and Answers

1. Difference between DDL, DML, DQL, and DCL

DDL (Data Definition Language):

- Commands: CREATE, ALTER, DROP, TRUNCATE
- Purpose: Define and modify database structure
- Example: `CREATE TABLE employees (...)`

DML (Data Manipulation Language):

- Commands: INSERT, UPDATE, DELETE
- Purpose: Manipulate data within tables
- Example: `INSERT INTO employees VALUES (...)`

DQL (Data Query Language):

- Commands: SELECT
- Purpose: Retrieve data from database
- Example: `SELECT * FROM employees WHERE salary > 50000`

DCL (Data Control Language):

- Commands: GRANT, REVOKE
- Purpose: Control access to database
- Example: `GRANT SELECT ON employees TO 'user'@'localhost'`

2. ACID Properties Example

sql

-- Atomicity: All operations in transaction succeed or all fail

START TRANSACTION;

UPDATE accounts SET balance = balance - 1000 WHERE account_id = 1;

UPDATE accounts SET balance = balance + 1000 WHERE account_id = 2;

COMMIT; *-- Both updates succeed or both fail*

-- Consistency: Database remains in valid state

-- Foreign key constraints ensure referential integrity

ALTER TABLE employees ADD CONSTRAINT fk_dept

FOREIGN KEY (dept_id) REFERENCES departments(dept_id);

-- Isolation: Transactions don't interfere with each other

-- Different isolation levels: READ UNCOMMITTED, READ COMMITTED, REPEATABLE READ, SERIALI

SET TRANSACTION ISOLATION LEVEL READ COMMITTED;

-- Durability: Committed changes persist

-- InnoDB storage engine ensures durability through write-ahead logging



3. Indexing Strategy Example

sql

```
-- When to create indexes:
-- 1. Primary keys (automatic)
-- 2. Foreign keys
CREATE INDEX idx_emp_dept ON employees(dept_id);

-- 3. Frequently searched columns
CREATE INDEX idx_emp_email ON employees(email);

-- 4. Columns used in ORDER BY
CREATE INDEX idx_emp_hire_date ON employees(hire_date);

-- 5. Composite indexes for multi-column searches
CREATE INDEX idx_emp_dept_salary ON employees(dept_id, salary);

-- When NOT to create indexes:
-- 1. Small tables (< 1000 rows)
-- 2. Columns that change frequently
-- 3. Tables with high INSERT/UPDATE/DELETE activity
```

Conclusion

This comprehensive guide covers all essential MySQL concepts using a practical employee management system example. The key areas covered include:

1. **Database Design:** Proper normalization and relationship modeling
2. **DDL Operations:** Creating and modifying database structures
3. **DML Operations:** Inserting, updating, and deleting data
4. **DQL Mastery:** Complex queries, joins, subqueries, and window functions

5. **Advanced Features:** Views, stored procedures, triggers, and JSON support
6. **Performance:** Indexing strategies and query optimization
7. **Security:** User management and access control
8. **Maintenance:** Backup, recovery, and monitoring

Next Steps for Students:

1. **Practice:** Set up the employee management database and run all examples
2. **Experiment:** Modify queries and observe the results
3. **Real Projects:** Apply these concepts to your own database projects
4. **Advanced Topics:** Explore MySQL 8.0 features like CTEs and window functions
5. **Performance Tuning:** Learn to use EXPLAIN and optimize slow queries

Additional Resources:

- MySQL Official Documentation
- MySQL Workbench for visual database design
- Performance monitoring tools like MySQL Enterprise Monitor
- Practice platforms like SQLBolt, W3Schools, and LeetCode SQL problems

Remember: The best way to learn MySQL is through hands-on practice. Start with simple queries and gradually work your way up to complex operations. Good luck with your MySQL journey!